


Paired system comparisons

Success-rate difference, A minus B · paired inference across 12 tasks



Bars: unadjusted 95% percentile intervals from 10,000 task-bootstrap resamples.
Holm p: exact task-label-swap tests; six contrasts. Intervals are not simultaneous.
Devin = Devin Desktop; Qwen = Qwen2.5-Coder-7B-Instruct.